



GROUP PROJECT – MIDTERM PRESENTATION

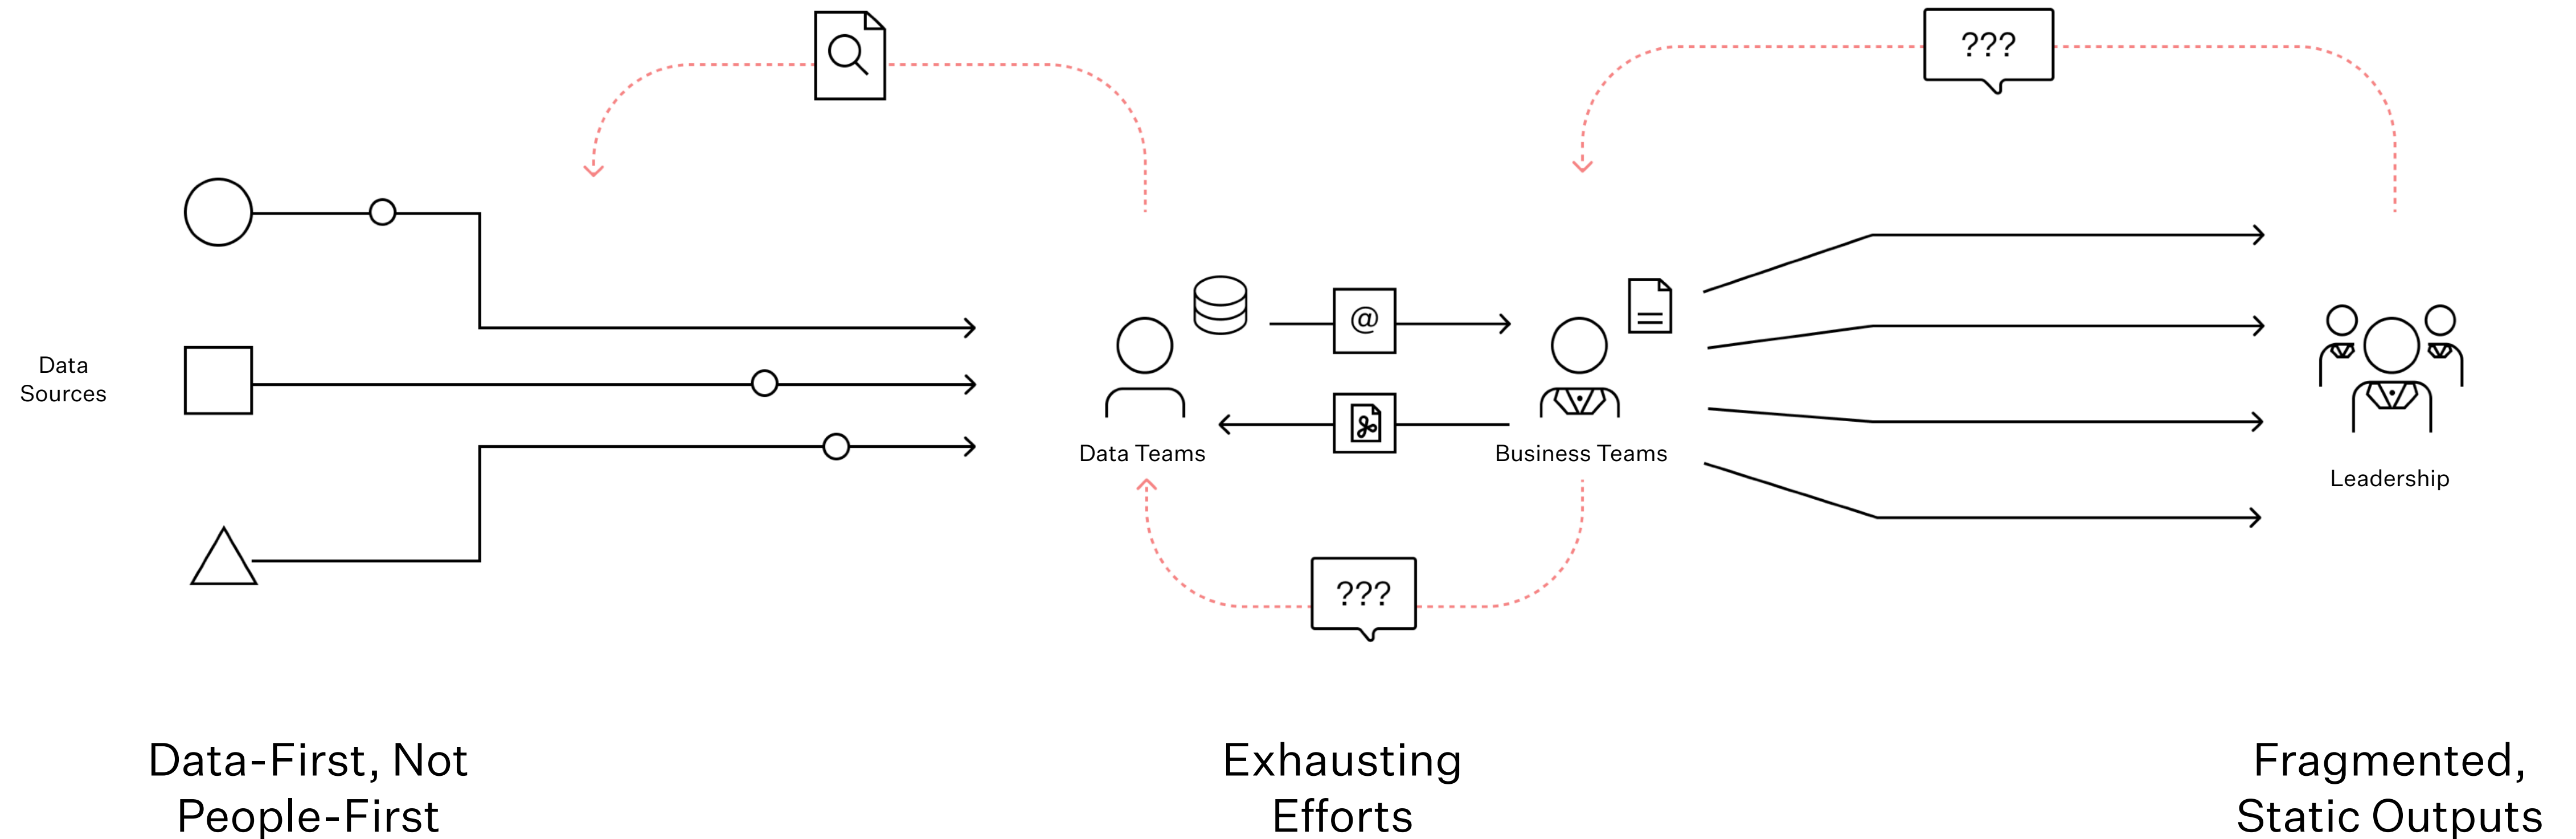
LumenAI: A Conversational Analytics Assistant for Intelligent Database Analysis

Mohamed Kassem, Erich Gunsenheimer | 27 MARCH 2025

MLOps & LLMOps: Production AI Systems - E115
SPRING 2025
Harvard Extension School



Overflowing Data, Elusive Clarity.





Solution

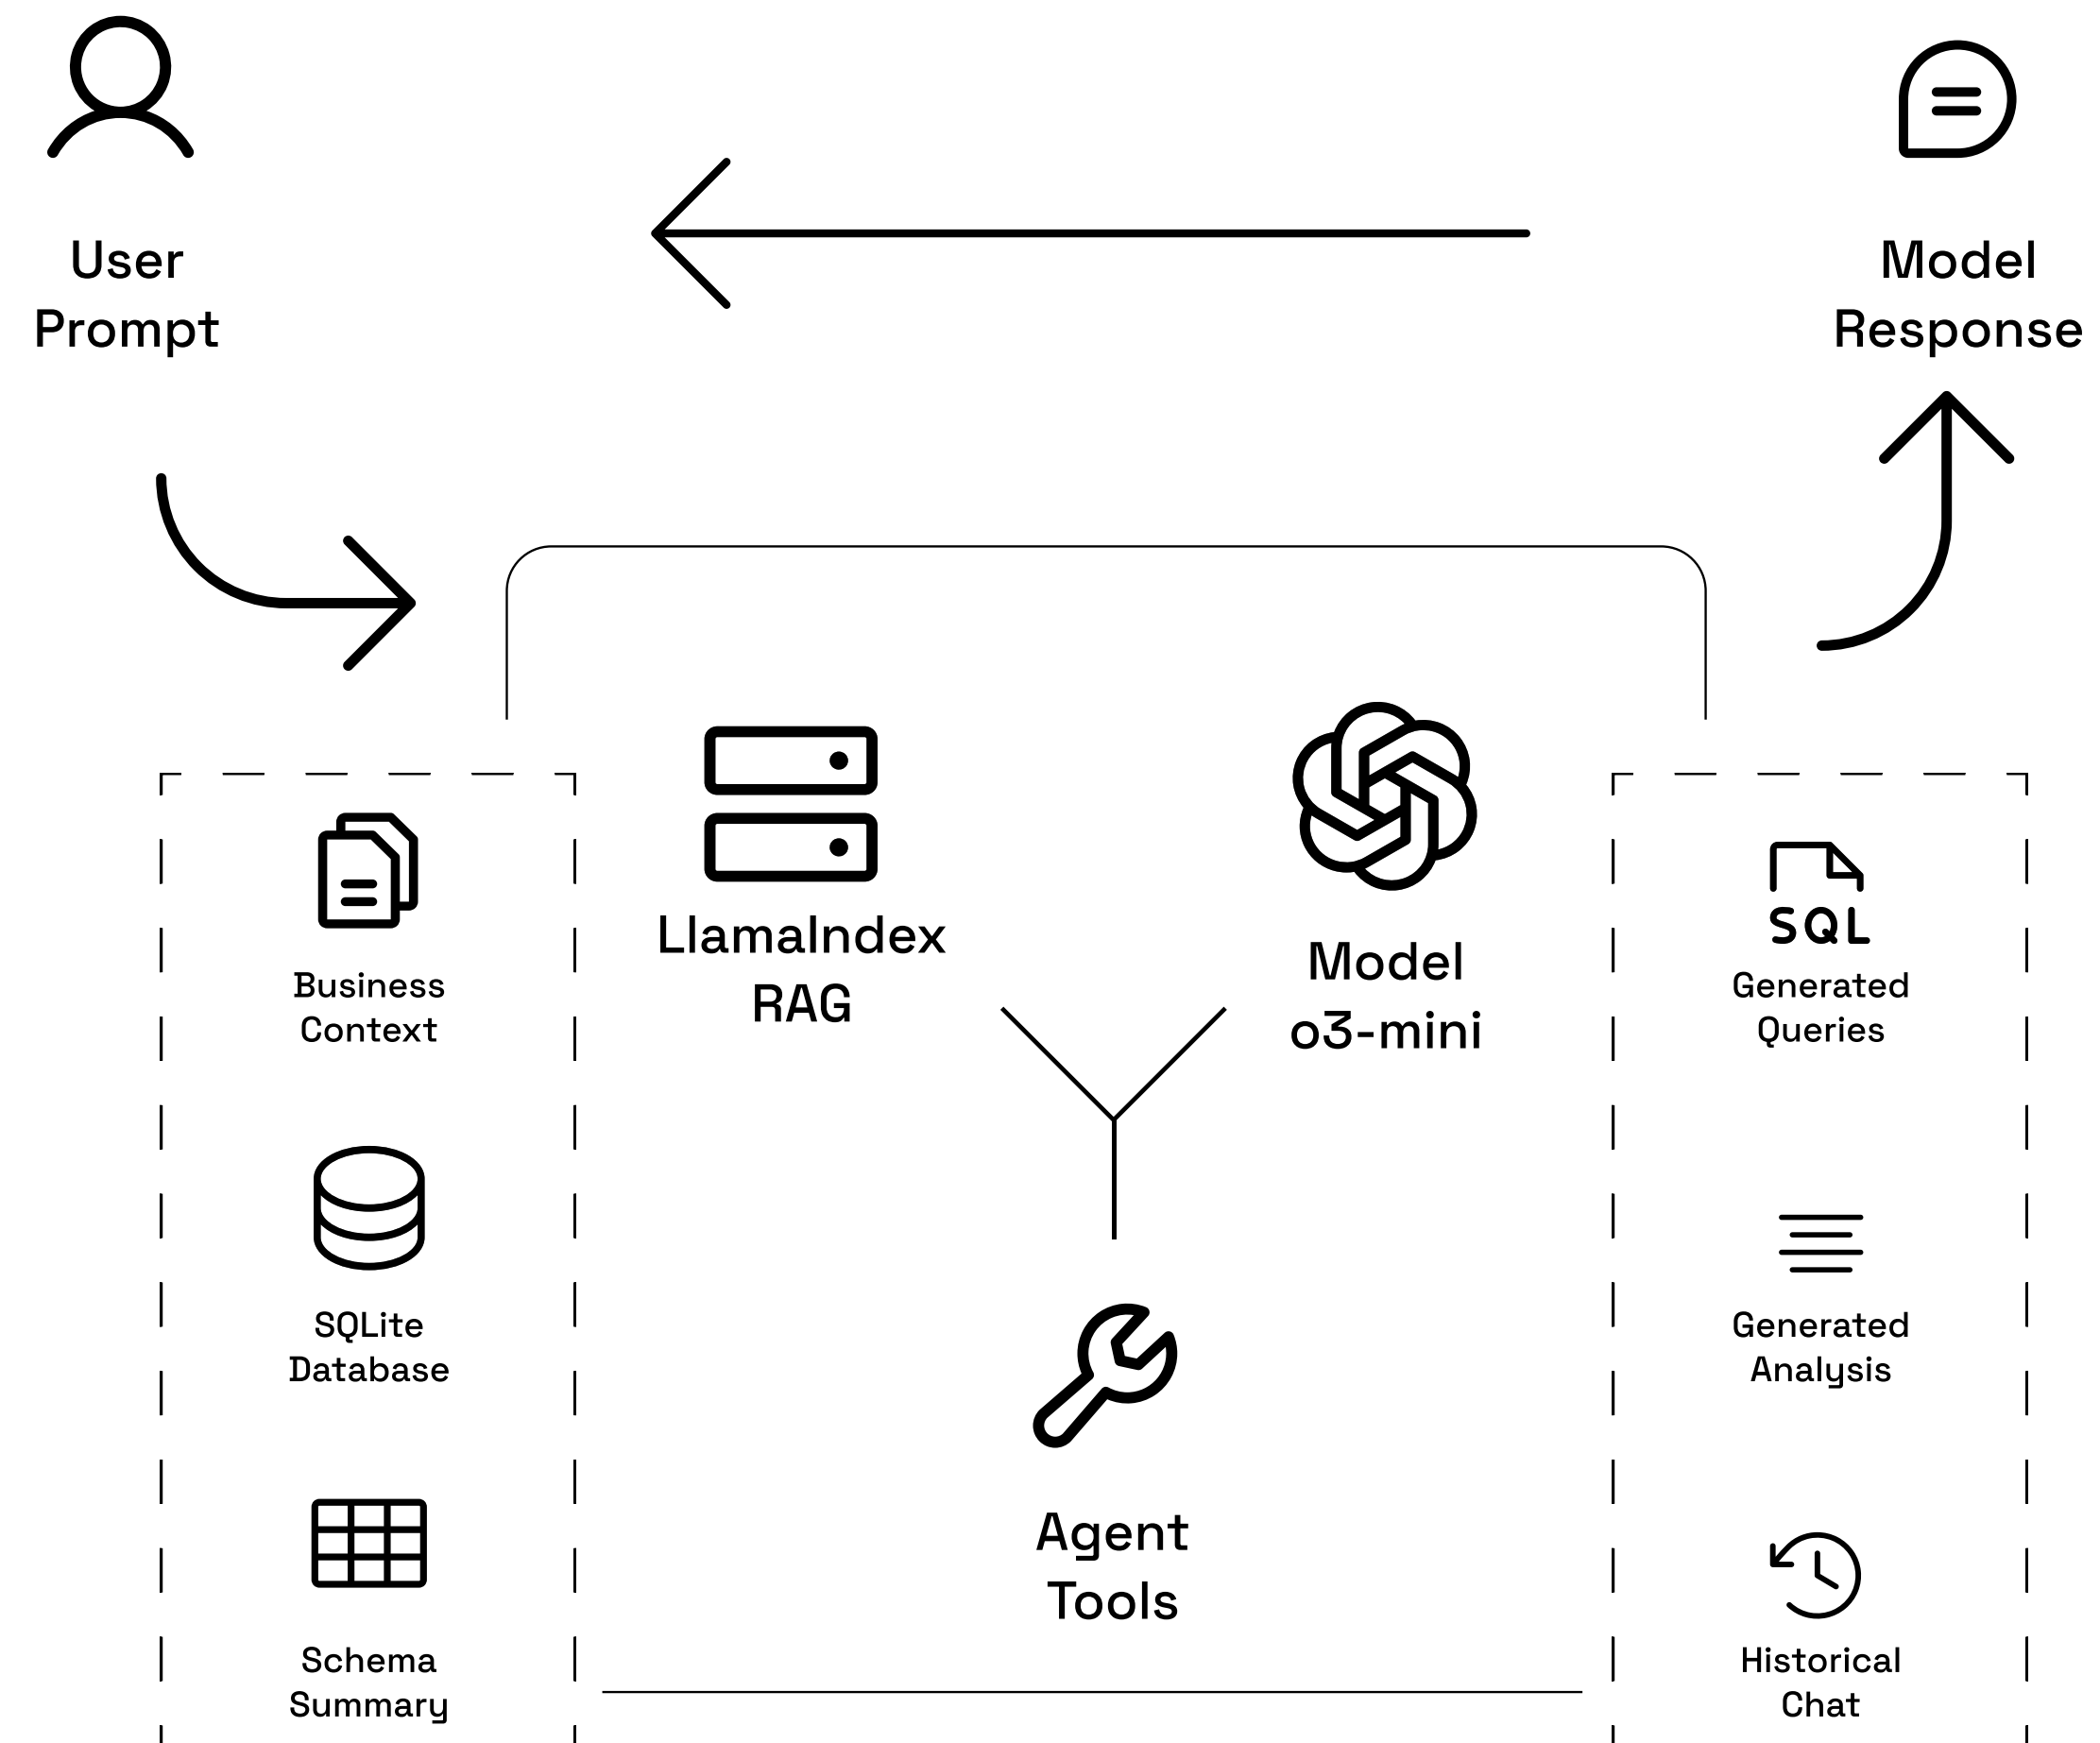
A context-rich ecosystem AI Agent that connects business users directly to data warehouse with flexible content management and reasoning

- Use Organization DB or adhoc datasets
LLM Summaries for schema and columns
- Incorporates previous presentations and reports
- Incorporates user context, role, and goals
- Incorporates historical context.
- Uses latest Integrated with validation rules, reasoning and tool selection.
- Tools include initiate SQL DB, Generate Queries, and specially-designed reporting prompts.



Technology

- Built on OpenAI o3-mini Model APIs
- Uses Embeddings from HuggingFace Embedding Model for vector embeddings
- Uses LlamaIndex for creating vector store and enabling semantic search over database schema
- Integrated with python functions
- Initiates SQLite DB through Python





Scalability & Efficiency

- Tested Multiple Model structure including locally run deepseek 3B Model and API-based OpenAI o3-mini model
- App is supported by dockershell, integrated with Git Repo, and runs on GCP



deepseek

OpenAI



Google Cloud Platform



GitHub



Scalability & Efficiency

Factors	#1	# 2	# 3	# 4	# 5
Model Used	Deepseek	o3-mini	o3-mini	o3-mini	o3-mini
Model Run	Locally	OpenAI	OpenAI	OpenAI	OpenAI
Context					
Schema	✓	✓	✓	LLM	LLM
Data	×	×	×	All — LLM	50 Rows — LLM
Statistics	×	✓	✓	LLM	LLM
Historical Prompts and Response	×	×	×	×	5 Records
Cost	Low	Low	Reasonable	Very High	Reasonable
Quality	Low	Low	Average	High	High



Future Development / Application

- Connect to User Friendly Interface
- Build ad-hoc document uploads
- Enhance the need-for-query validation workflow
- Enhance Error-Handling Workflow and Reasoning
- Enhance Report Prompting
- Optimize running cost



Future Development / Growth

- Introduce Domain-Refined Versions
- Optimize locally hosted models for more secured environments
- Integrate with Data Management systems for deeper context on lineage, usage, sources, and others
- Enhance KPI Suggestion
- Build a unified pool of queries/governance on an organization level



Thank You